

MACHINE LEARNING LAB

EXERCISE 2

Aim :

Use the house_pred.csv file to build a multiple linear regression model. sklearn shall be used to fit the model.

Perform necessary preprocessing and check for outliers and multi-collinearity. Apply the same set of preprocessing to the test.csv and use the data to predict the house price. The evaluation criteria will be Root Mean Squared Error

Algorithm :

1. Data Preprocessing and Exploration:

- Load and preprocess `house_pred.csv` , handling missing values and outliers.
- Explore data with EDA to understand features and relationships.

2. Model Building and Evaluation:

- Split data into train and validation sets.
- Use sklearn to build a multiple linear regression model.
- Train the model and evaluate with RMSE.

3. Prediction on Test Data:

- Apply preprocessing steps to `test.csv` .
- Use trained model to predict house prices.
- Evaluate predictions with RMSE.

Code and Output :

In [1]:

```
pip install Numpy==1.23.5
```

Requirement already satisfied: Numpy==1.23.5 in c:\users\teju\anaconda3\lib\site-packages (1.23.5)

Note: you may need to restart the kernel to use updated packages.

In [2]:

```
import pandas as pd
from statsmodels.stats.outliers_influence import variance_inflation_factor as VIF
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error
import math
```

C:\Users\TEJU\anaconda3\lib\site-packages\scipy__init__.py:146: UserWarning: A NumPy version >=1.16.5 and <1.23.0 is required for this version of SciPy (detected version 1.23.5
warnings.warn(f"A NumPy version >={np_minversion} and <{np_maxversion}")

```
In [3]: data=pd.read_csv(r"C:\Users\TEJU\Downloads\house_pred (1).csv")
data.head()
```

```
Out[3]:
```

	Id	MSSubClass	MSZoning	LotFrontage	LotArea	Street	Alley	LotShape	LandContour	Utilities
0	1	60	RL	65.0	8450	Pave	NaN	Reg	Lvl	AllPub
1	2	20	RL	80.0	9600	Pave	NaN	Reg	Lvl	AllPub
2	3	60	RL	68.0	11250	Pave	NaN	IR1	Lvl	AllPub
3	4	70	RL	60.0	9550	Pave	NaN	IR1	Lvl	AllPub
4	5	60	RL	84.0	14260	Pave	NaN	IR1	Lvl	AllPub

5 rows × 81 columns



```
In [4]: data.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1460 entries, 0 to 1459
Data columns (total 81 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Id                    1460 non-null   int64
1   MSSubClass            1460 non-null   int64
2   MSZoning              1460 non-null   object
3   LotFrontage          1201 non-null   float64
4   LotArea               1460 non-null   int64
5   Street               1460 non-null   object
6   Alley                91 non-null     object
7   LotShape             1460 non-null   object
8   LandContour          1460 non-null   object
9   Utilities            1460 non-null   object
10  LotConfig            1460 non-null   object
11  LandSlope            1460 non-null   object
12  Neighborhood          1460 non-null   object
13  Condition1           1460 non-null   object
14  Condition2           1460 non-null   object
15  BldgType             1460 non-null   object
16  HouseStyle           1460 non-null   object
17  OverallQual          1460 non-null   int64
18  OverallCond          1460 non-null   int64
19  YearBuilt            1460 non-null   int64
20  YearRemodAdd         1460 non-null   int64
21  RoofStyle            1460 non-null   object
22  RoofMatl            1460 non-null   object
23  Exterior1st          1460 non-null   object
24  Exterior2nd          1460 non-null   object
25  MasVnrType           1452 non-null   object
26  MasVnrArea           1452 non-null   float64
27  ExterQual            1460 non-null   object
28  ExterCond            1460 non-null   object
29  Foundation           1460 non-null   object
30  BsmtQual             1423 non-null   object
31  BsmtCond            1423 non-null   object
32  BsmtExposure         1422 non-null   object
33  BsmtFinType1         1423 non-null   object
34  BsmtFinSF1          1460 non-null   int64
```

```

35 BsmtFinType2    1422 non-null    object
36 BsmtFinSF2      1460 non-null    int64
37 BsmtUnfSF       1460 non-null    int64
38 TotalBsmtSF     1460 non-null    int64
39 Heating         1460 non-null    object
40 HeatingQC       1460 non-null    object
41 CentralAir      1460 non-null    object
42 Electrical      1459 non-null    object
43 1stFlrSF        1460 non-null    int64
44 2ndFlrSF        1460 non-null    int64
45 LowQualFinSF    1460 non-null    int64
46 GrLivArea       1460 non-null    int64
47 BsmtFullBath    1460 non-null    int64
48 BsmtHalfBath    1460 non-null    int64
49 FullBath        1460 non-null    int64
50 HalfBath        1460 non-null    int64
51 BedroomAbvGr   1460 non-null    int64
52 KitchenAbvGr   1460 non-null    int64
53 KitchenQual     1460 non-null    object
54 TotRmsAbvGrd   1460 non-null    int64
55 Functional      1460 non-null    object
56 Fireplaces      1460 non-null    int64
57 FireplaceQu     770 non-null     object
58 GarageType      1379 non-null    object
59 GarageYrBlt     1379 non-null    float64
60 GarageFinish    1379 non-null    object
61 GarageCars      1460 non-null    int64
62 GarageArea      1460 non-null    int64
63 GarageQual      1379 non-null    object
64 GarageCond      1379 non-null    object
65 PavedDrive      1460 non-null    object
66 WoodDeckSF      1460 non-null    int64
67 OpenPorchSF     1460 non-null    int64
68 EnclosedPorch   1460 non-null    int64
69 3SsnPorch       1460 non-null    int64
70 ScreenPorch     1460 non-null    int64
71 PoolArea        1460 non-null    int64
72 PoolQC          7 non-null       object
73 Fence           281 non-null     object
74 MiscFeature     54 non-null      object
75 MiscVal         1460 non-null    int64
76 MoSold          1460 non-null    int64
77 YrSold          1460 non-null    int64
78 SaleType        1460 non-null    object
79 SaleCondition   1460 non-null    object
80 SalePrice       1460 non-null    int64

```

dtypes: float64(3), int64(35), object(43)

memory usage: 924.0+ KB

In [5]:

```

def ThresholdandND_columnRemoval(df):
    N = len(df)
    columns = df.columns
    for col in columns:
        if (len(df[col].unique()) == 1):
            df = df.drop([col],axis=1)
            continue
        notnull = df[col].isnull().sum()
        ratio = notnull / N
        if(ratio >= 0.30):
            df = df.drop([col],axis=1)
    return df

def Handling_NullValues(df):
    columns = df.columns

```

```

for col in columns:
    typeCol = str(df[col].dtype)
    if typeCol == 'object':
        df = df[df[col].notna()]
    else:
        mean = df[col].mean()
        median = df[col].median()
        standard_deviation = df[col].std()
        pmc = (3 * (mean - median)) / standard_deviation
        if pmc >= 0.4 or pmc <= -0.4:
            df[col] = df[col].fillna(median)
        else:
            df[col] = df[col].fillna(mean)
return df

def OneHotEncoding_objects(df):
    columns = df.columns
    for col in columns:
        typeCol = str(df[col].dtype)
        if typeCol == 'object':
            enc = pd.get_dummies(df[col])
            encCol = enc.columns
            newColumns = {}
            for i in range(0, len(encCol)):
                newColumns[encCol[i]] = col + encCol[i]
            enc.rename(columns = newColumns, inplace = True)
            df = df.join(enc)
            df = df.drop([col], axis=1)
    return df

def IQR_Removal(df):
    columns = df.columns
    for col in columns:
        if col == 'SalePrice':
            continue
        typeCol = str(df[col].dtype)
        if typeCol != 'object':
            Q1 = df[col].quantile(0.25)
            Q3 = df[col].quantile(0.75)
            iqr = Q3 - Q1
            df = df[(df[col] >= Q1 - 1.5*iqr) & (df[col] <= Q3 + 1.5*iqr)]
    return df

```

In [6]:

```
df = OneHotEncoding_objects(IQR_Removal(Handling_NullValues(ThresholdandND_columnRem
df
```

Out[6]:

	Id	MSSubClass	LotFrontage	LotArea	OverallQual	OverallCond	YearBuilt	YearRemodAdd
0	1	60	65.0	8450	7	5	2003	2003
2	3	60	68.0	11250	7	5	2001	2002
4	5	60	84.0	14260	8	5	2000	2000
6	7	20	75.0	10084	8	5	2004	2005
10	11	20	70.0	11200	5	5	1965	1965
...
1447	1448	60	80.0	10000	8	5	1995	1996
1448	1449	50	70.0	11767	4	7	1910	2000

	Id	MSSubClass	LotFrontage	LotArea	OverallQual	OverallCond	YearBuilt	YearRemodAdd
1451	1452	20	78.0	9262	8	5	2008	2009
1454	1455	20	62.0	7500	7	5	2004	2005
1455	1456	60	62.0	7917	6	5	1999	2000

535 rows × 211 columns

In [7]:

```
def VIF_Filter(df,dfTest):
    xCol = list(set(list(df.columns)) & set(list(dfTest.columns)))
    xCol.remove('Id')
    while(1):
        finished = True
        xVal = df[xCol]
        xVal['intercept'] = 1
        vif = pd.DataFrame()
        vif['variable'] = xVal.columns
        vif['vif'] = [VIF(xVal.values,i) for i in range(xVal.shape[1])]
        for i in range(0,len(vif)):
            var = str(vif.iloc[i,0])
            val = str(vif.iloc[i,1])
            if(var == 'intercept'):
                continue
            elif(val == 'inf'):
                xCol.remove(var)
                finished = False
                break
            else:
                val = float(val)
                if val > 3:
                    xCol.remove(var)
                    finished = False
                    break
        if finished == True:
            return xCol, vif
```

In [10]:

```
testdata = pd.read_csv(r"C:\Users\TEJU\Downloads\test (1).csv")
testdf = OneHotEncoding_objects(IQR_Removal(Handling_NullValues(ThresholdandND_column
columns, vif = VIF_Filter(df,testdf)
```

C:\Users\TEJU\AppData\Local\Temp\ipykernel_16600\2458712660.py:7: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
xVal['intercept'] = 1
```

C:\Users\TEJU\anaconda3\lib\site-packages\statsmodels\stats\outliers_influence.py:193: RuntimeWarning: divide by zero encountered in double_scalars

```
vif = 1. / (1. - r_squared_i)
```

C:\Users\TEJU\anaconda3\lib\site-packages\statsmodels\regression\linear_model.py:1715: RuntimeWarning: divide by zero encountered in double_scalars


```
return 1 - self.ssr/self.centered_tss
```

C:\Users\TEJU\anaconda3\lib\site-packages\statsmodels\regression\linear_model.py:1715: RuntimeWarning: invalid value encountered in double_scalars

```
return 1 - self.ssr/self.centered_tss
```

```
In [12]: corr = df[columns].corr()
f, ax = plt.subplots(figsize=(12, 10))
mask = np.triu(np.ones_like(corr, dtype=bool))
cmap = sns.diverging_palette(230, 20, as_cmap=True)
sns.heatmap(corr, mask = mask, cmap=cmap)
```

Out[12]: <AxesSubplot:>



```
In [14]: y = df['SalePrice']
x = df[columns]
```

```
reg = LinearRegression()
reg.fit(x,y)
```

Out[14]:

LinearRegression ⓘ ⓘ

LinearRegression()

In [15]:

```
reg.intercept_, reg.coef_
```

Out[15]:

```
(-1594359.180326281,
 array([ 1.48196131e+04,  1.79923581e+03, -2.01900480e+04,  8.29220098e+03,
        -3.11602711e+04, -2.13025379e+04,  9.85542569e+03,  1.01402363e+04,
        -4.80839507e+01, -6.51601042e+03,  4.32819720e+03, -1.10285328e-08,
         1.54644927e+03,  1.10831132e+04,  6.89396984e-10, -1.80896810e+03,
         2.60902265e+03,  1.70985004e-10,  9.34399022e+03, -1.51054025e+04,
        -6.89746609e+03, -3.63797881e-12, -1.49015826e+03,  1.58803862e+04,
         2.64284679e+04, -3.24440391e+04, -6.77657514e+03,  2.90603693e+01,
         3.06347254e+02,  6.18456397e-11, -3.81987775e-11, -3.23539744e+04,
         2.95778294e+03,  6.48982182e+02,  1.86446414e-11, -5.02389241e-11,
         7.86574452e+02, -5.45696821e-12, -1.59962832e+02,  6.81430689e+03,
         1.49720747e+04, -3.79134049e+04, -4.44976327e+02,  9.50544655e+03,
         4.79540257e+04,  8.62056422e+03,  2.85956868e+04,  3.30992550e+04,
        -1.31876732e-11,  0.00000000e+00, -6.49579038e+03,  2.00776318e+04,
         3.63797881e-12,  2.90029219e+04, -2.35091579e+04, -8.45164260e+02,
         1.61967148e+02,  1.30491980e+04,  5.67350115e+03,  1.01211579e+03,
         1.12406842e+03,  1.58823688e+04, -1.38634698e+04,  3.17295794e+04,
         2.21720800e+04,  2.02833624e+03,  9.75666077e+03,  1.67018922e+04,
        -4.51834828e+03, -3.87905174e+02, -8.93944758e+03, -1.16442061e+04,
        -3.03759292e+03,  4.98591671e+03,  4.67010521e+02,  1.10556570e+04,
         1.59655180e+04, -1.45586274e+04,  3.33199304e+01,  5.35591104e+03,
         2.33390729e+01]))
```

In [17]:

```
x_test = testdf.columns]
y_forTest = reg.predict(x_test)
y_forTest
```

Out[17]:

```
array([176378.22648779, 208983.52778748, 290000.          , 259602.35385925,
       196874.42652567, 110157.92856467, 142053.28850765, 147620.23703439,
       190748.79104513, 225113.75690499, 174088.87649298, 206004.63007601,
       203899.99223414, 221843.22098781, 178814.32583258, 151970.68978991,
       191895.4718545 , 249476.85597501, 158440.66623878, 131101.521714 ,
       238783.07427756, 165335.360849 , 211234.4214451 , 170922.61185717,
       230529.38276664, 164397.9115744 , 188796.53885744, 196219.72656122,
       228279.65302757, 231189.58306462, 192178.72629544, 196539.36220642,
       174165.52529681, 173215.57956365, 206689.93686828, 115634.04844513,
       174531.9662555 ])
```

In [18]:

```
y_pred = reg.predict(x)
math.sqrt(mean_squared_error(y,y_pred))
```

Out[18]:

```
23373.573363257263
```

In [19]:

```
reg.score(x,y)
```

Out[19]:

```
0.8460571194156481
```

Results :

Therefore, we were successfully able to build the multiple linear regression model and use the train data to predict the house prices in the test data.