

MACHINE LEARNING LAB

EXERCISE 5

Aim :

Use the teleco-customer-churn dataset for the following:

1. Perform the necessary pre-processings.
2. Apply all the classification algorithms (KNN, Logistic Regression, Naive Bayes, Decision Trees, SVM) on this dataset and print the accuracies.
3. Find which algorithm gave the best accuracy.
4. Provide a justification as to why that algorithm provided the best accuracy

Algorithm :

1. We load and preprocess the data by removing the unnecessary features, hot encoding.
1. We then split the data into X and Y and then scale the data
1. One by one, we apply the classification algorithms, KNN, Logistic Regression, Naive Bayes, Decision Trees and SVM on the dataset
1. We also print the performance metrics using each of the algorithms.
1. We then try to analyse why a particular algorithm would have given the highest accuracy and justify.

Code and Output :

```
In [72]: import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
```

```
In [73]: df=pd.read_csv(r"C:\Users\TEJU\Downloads\Telco-Customer-Churn.csv")
```

```
In [74]: df.head()
```

```
Out[74]:
```

	customerID	gender	SeniorCitizen	Partner	Dependents	tenure	PhoneService	MultipleLines	Int
0	7590-VHVEG	Female	0	Yes	No	1	No	No phone service	
1	5575-GNVDE	Male	0	No	No	34	Yes	No	

	customerID	gender	SeniorCitizen	Partner	Dependents	tenure	PhoneService	MultipleLines	Int
2	3668-QPYBK	Male	0	No	No	2	Yes	No	
3	7795-CFOCW	Male	0	No	No	45	No	No phone service	
4	9237-HQITU	Female	0	No	No	2	Yes	No	

5 rows × 21 columns

In [75]:

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 7043 entries, 0 to 7042
Data columns (total 21 columns):
#   Column                Non-Null Count  Dtype
---  -
0   customerID            7043 non-null   object
1   gender                 7043 non-null   object
2   SeniorCitizen          7043 non-null   int64
3   Partner                7043 non-null   object
4   Dependents             7043 non-null   object
5   tenure                 7043 non-null   int64
6   PhoneService           7043 non-null   object
7   MultipleLines          7043 non-null   object
8   InternetService        7043 non-null   object
9   OnlineSecurity         7043 non-null   object
10  OnlineBackup           7043 non-null   object
11  DeviceProtection       7043 non-null   object
12  TechSupport            7043 non-null   object
13  StreamingTV            7043 non-null   object
14  StreamingMovies        7043 non-null   object
15  Contract               7043 non-null   object
16  PaperlessBilling       7043 non-null   object
17  PaymentMethod          7043 non-null   object
18  MonthlyCharges         7043 non-null   float64
19  TotalCharges           7043 non-null   object
20  Churn                  7043 non-null   object
dtypes: float64(1), int64(2), object(18)
memory usage: 1.1+ MB
```

I notice that a particular numerical column TotalCharges is given as an object type instead of int/float. So we change that first

In [76]:

```
df['TotalCharges'].replace(" ",0,inplace=True)
df['TotalCharges']=df['TotalCharges'].astype('float64')
```

We drop the customerID column as it is not required

In [77]:

```
df=df.drop('customerID',axis=1)
```

We give meaning to the 0's and 1's in the senior citizen column by mapping to No's and Yes's

```
In [78]: df['SeniorCitizen']=df['SeniorCitizen'].map({0:'No',1:'Yes'})
```

We separate the columns having numbers and columns having words separately

```
In [79]: num_features=df.select_dtypes(include='number')
cat_features=df.select_dtypes(exclude='number')
```

Let us now perform hot encoding on the cat features and then calculate correlation matrix

```
In [80]: cat_features_encoded = pd.get_dummies(data=cat_features, dtype=int)
churn_corr = cat_features_encoded.corr()['Churn_Yes'].drop(['Churn_Yes', 'Churn_No'])
```

```
In [81]: df_final=pd.get_dummies(data=df,drop_first=True,dtype=int)
```

We drop the features which are unrelated to the target variable churn


```
In [82]: df_final = df_final.drop(['gender_Male', 'PhoneService_Yes',
                                   'MultipleLines_No phone service',
                                   'MultipleLines_Yes'], axis=1)
```

```
In [83]: df_final.head()
```

```
Out[83]:
```

	tenure	MonthlyCharges	TotalCharges	SeniorCitizen_Yes	Partner_Yes	Dependents_Yes	InternetService
0	1	29.85	29.85	0	1	0	
1	34	56.95	1889.50	0	0	0	
2	2	53.85	108.15	0	0	0	
3	45	42.30	1840.75	0	0	0	
4	2	70.70	151.65	0	0	0	

5 rows × 27 columns



```
In [84]: df_final['Churn_Yes'].value_counts()
```

```
Out[84]: 0    5174
         1    1869
         Name: Churn_Yes, dtype: int64
```

We notice that there is an imbalance of 0's and 1's because of which the model will tend to give 0 as the answer because of its high number, we can use some techniques to adjust that by adding random values to increase number of 1's like SMOTE but refraining from doing the same for this lab

Let us now split into X and Y to start creating our ML model

```
In [85]: X=df_final.drop('Churn_Yes',axis=1)
         y=df_final['Churn_Yes']
```

```
In [87]: from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_stat
```

Let us scale the data now

```
In [88]: from sklearn.preprocessing import StandardScaler
scaler = StandardScaler()
X_train = scaler.fit_transform(X_train)
X_test = scaler.transform(X_test)
```

KNN Classifier

```
In [89]: from sklearn.neighbors import KNeighborsClassifier
```

```
In [107... knn_model = KNeighborsClassifier(n_jobs=-1, n_neighbors=5)
knn_model.fit(X_train, y_train)
```

```
Out[107... KNeighborsClassifier(n_jobs=-1)
```

```
In [108... y_pred = knn_model.predict(X_test)
```

```
In [109... from sklearn.metrics import confusion_matrix, accuracy_score, classification_report
```

```
In [110... confusion_matrix(y_test, y_pred)
```

```
Out[110... array([[872, 161],
       [184, 192]], dtype=int64)
```

```
In [111... accuracy_score(y_test, y_pred)
```

```
Out[111... 0.7551454932576295
```

```
In [162... class_report = classification_report(y_test, y_pred)
print(class_report)
```

	precision	recall	f1-score	support
0	0.83	0.88	0.85	1033
1	0.60	0.52	0.56	376
accuracy			0.78	1409
macro avg	0.72	0.70	0.71	1409
weighted avg	0.77	0.78	0.78	1409

=> KNN Classifier - Accuracy of 75.5%

Logistic Regression

```
In [113... from sklearn.linear_model import LogisticRegression
reg = LogisticRegression()
reg.fit(X_train,y_train)
```

```
Out[113... LogisticRegression()
```

```
In [114... y_pred = reg.predict(X_test)
```

```
In [115... confusion_matrix(y_test,y_pred)
```

```
Out[115... array([[917, 116],
        [171, 205]], dtype=int64)
```

```
In [116... accuracy_score(y_test,y_pred)
```

```
Out[116... 0.7963094393186657
```

```
In [117... class_report=classification_report(y_test,y_pred)
print(class_report)
```

	precision	recall	f1-score	support
0	0.84	0.89	0.86	1033
1	0.64	0.55	0.59	376
accuracy			0.80	1409
macro avg	0.74	0.72	0.73	1409
weighted avg	0.79	0.80	0.79	1409

=> Logistic Regression - Accuracy of 79.6%

Naive Bayes

```
In [118... from sklearn.naive_bayes import GaussianNB
classifier = GaussianNB()
classifier.fit(X_train,y_train)
```

```
Out[118... GaussianNB()
```

```
In [119... y_pred = classifier.predict(X_test)
```

```
In [120... confusion_matrix(y_test,y_pred)
```

```
Out[120... array([[597, 436],
        [ 52, 324]], dtype=int64)
```

```
In [121... accuracy_score(y_test,y_pred)
```

```
Out[121... 0.6536550745209369
```

In [123...

```
class_report=classification_report(y_test,y_pred)
print(class_report)
```

	precision	recall	f1-score	support
0	0.92	0.58	0.71	1033
1	0.43	0.86	0.57	376
accuracy			0.65	1409
macro avg	0.67	0.72	0.64	1409
weighted avg	0.79	0.65	0.67	1409

=> Naive Bayes - Accuracy of 65.3%

Decision Trees

In [124...

```
from sklearn.tree import DecisionTreeClassifier
tree = DecisionTreeClassifier(criterion='entropy')
tree.fit(X_train,y_train)
```

Out[124...

```
DecisionTreeClassifier(criterion='entropy')
```

In [125...

```
y_pred = tree.predict(X_test)
```

In [126...

```
confusion_matrix(y_test,y_pred)
```

Out[126...

```
array([[859, 174],
       [190, 186]], dtype=int64)
```

In [127...

```
accuracy_score(y_test,y_pred)
```

Out[127...

```
0.7416607523066004
```

In [129...

```
class_report=classification_report(y_test,y_pred)
print(class_report)
```

	precision	recall	f1-score	support
0	0.82	0.83	0.83	1033
1	0.52	0.49	0.51	376
accuracy			0.74	1409
macro avg	0.67	0.66	0.67	1409
weighted avg	0.74	0.74	0.74	1409

In [130...

```
tree2=DecisionTreeClassifier()
tree2.fit(X_train,y_train)
```

Out[130...

```
DecisionTreeClassifier()
```

In [131...

```
y_pred=tree2.predict(X_test)
```

```
In [132... confusion_matrix(y_test,y_pred)
```

```
Out[132... array([[840, 193],
        [188, 188]], dtype=int64)
```

```
In [133... accuracy_score(y_test,y_pred)
```

```
Out[133... 0.7295954577714692
```

```
In [134... class_report = classification_report(y_test, y_pred)
print(class_report)
```

	precision	recall	f1-score	support
0	0.82	0.81	0.82	1033
1	0.49	0.50	0.50	376
accuracy			0.73	1409
macro avg	0.66	0.66	0.66	1409
weighted avg	0.73	0.73	0.73	1409

I tried with both criterions of decision tree - Gini Impurity and entropy.

Entropy gave an accuracy higher than Gini

=> **Decision Tree - Accuracy of 74.1%**

Support Vector Machines

```
In [156... from sklearn.svm import SVC
clf = SVC(kernel='linear')
```

```
In [157... clf.fit(X_train,y_train)
```

```
Out[157... SVC(kernel='linear')
```

```
In [158... y_pred=clf.predict(X_test)
```

```
In [159... confusion_matrix(y_test,y_pred)
```

```
Out[159... array([[905, 128],
        [180, 196]], dtype=int64)
```

```
In [160... accuracy_score(y_test,y_pred)
```

```
Out[160... 0.7814052519517388
```

```
In [161... class_report = classification_report(y_test, y_pred)
print(class_report)
```

	precision	recall	f1-score	support
0	0.83	0.88	0.85	1033
1	0.60	0.52	0.56	376
accuracy			0.78	1409
macro avg	0.72	0.70	0.71	1409
weighted avg	0.77	0.78	0.78	1409

Tried different kernels like linear, poly, rbf and sigmoid

Got the highest accuracy for linear

=> SVM - Accuracy of 78.1%

Justification file :

Algorithm with best accuracy :

All five algorithms provided accuracy around the same range but the algorithm with the highest accuracy is Logistic Regression with an accuracy of **79.6%**.

The reason why accuracy isn't very high is because of the nature of the dataset where in the target variable 'churn' had high number of No's and less number of Yes's (0's and 1's) because of which there is an imbalance.

The reason why logistic regression has the highest accuracy is as follows :

- Let us notice the F1 score closely. F1 score is the measure of the harmonic mean of precision and recall. Commonly used as an evaluation metric in binary and multi-class classification, the F1 score integrates precision and recall into a single metric to gain a better understanding of model performance. It is 0.86 and 0.59 which signifies that the model can effectively identify positive cases while minimising false positives and false negatives.
- Precision is also higher than the ones achieved in other algorithms with 0.84 and 0.64
- Dataset is fairly linear allowing better capturing in this model

Result :

Therefore, we were successfully able to apply all the classification algorithms (KNN, Logistic Regression, Naive Bayes, Decision Trees, SVM) on this dataset and print the accuracies. Additionally, we were able to justify reason for highest accuracy algorithm.